

FYS-STK3155 Project 1

Lasse Pladsen, Parham Qanbari, & Sander V. Vattøy

September 30, 2023

Abstract

We have studied three different regression methods, ordinary least squares, ridge, and lasso, and what affects their learned prediction accuracy. We found We have done a bias-variance analysis on OLS using a bootstrap resampling technique and found that We then study the cross-validation resampling method and analyze all three regression methods Finally we applied our models to real world topographic data and we saw that

I Introduction

In today's modern world machine learning has emerged as a revolutionary technology. Even if just superficially, machine learning has become well known around the world. By allowing computers to learn from supplied data we create powerful tools with incredible real world applications for analysis and prediction.

In this project we will explore and study low-level machine learning methods. Our focus will be on understanding and optimizing the potential of data regression methods such as ordinary least squares (OLS), ridge regression, and lasso regression. We will look into the factors that influence their accuracy and predictive capabilities, employing a variety of strategies to minimize errors and improve performance. Finally, we will apply our finely-tuned algorithms to analyze real-world topographic data.

II Theory

A Design matrix for linear regression

We create a so-called design matrix \mathbf{X} for the regression methods from two input variables \mathbf{x} and \mathbf{y} . This matrix is used to create our linear regression models. Each row in \mathbf{X} represents polynomial variables from one data sample. We choose a max polynomial

degree p with n data samples such that $\mathbf{X} \in \mathbb{R}^{n \times l}$ where

$$l = \text{floor}\left[\frac{1}{2}(p+1)(p+2)\right]$$

and the design matrix is then given by

$$\mathbf{x} = \begin{bmatrix} 1 & x_0 & y_0 & x_0^2 & x_0 y_0 & y_0^2 & \dots & x_0^{p-1} & \dots & y_0^{p-1} \\ 1 & x_1 & y_1 & x_1^2 & x_1 y_1 & y_1^2 & \dots & x_1^{p-1} & \dots & y_1^{p-1} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_{n-1} & y_{n-1} & x_{n-1}^2 & x_{n-1} y_{n-1} & y_{n-1}^2 & \dots & x_{n-1}^{p-1} & \dots & y_{n-1}^{p-1} \end{bmatrix} \quad (1)$$

We then make our linear regression predictions by creating the polynomials from \mathbf{X} and the polynomial coefficients $\boldsymbol{\beta} \in \mathbb{R}^l$ as follows

$$\tilde{\mathbf{z}} = \mathbf{X}\boldsymbol{\beta} \quad (2)$$

B Regression methods

The different methods have different ways of calculating the optimal $\boldsymbol{\beta}$ -coefficients which we call $\hat{\boldsymbol{\beta}}$.

B.1 Ordinary least squares

$$\hat{\boldsymbol{\beta}}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{z} \quad (3)$$

B.2 Ridge

$$\hat{\boldsymbol{\beta}}_{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{z} \quad (4)$$

where \mathbf{I} is the identity matrix, and λ is a non-negative regularization parameter.

B.3 Lasso

For lasso regression we do not have an analytical expression, but we will use the functionalities of the machine learning python module **SciKit-Learn**.

B.4 Error and coefficient of determination

To calculate our regression models' prediction error from the actual data we will be using the mean square error (MSE) defined by

$$MSE(\mathbf{z}, \tilde{\mathbf{z}}_i) = \frac{1}{n} \sum_{i=0}^{n-1} (z_i - \tilde{z}_i)^2, \quad (5)$$

We calculate the coefficient of determination for our model, called the R^2 , given as

$$R^2(\mathbf{z}, \tilde{\mathbf{z}}_i) = 1 - \frac{\sum_{i=0}^{n-1} (z_i - \tilde{z}_i)^2}{\sum_{i=0}^{n-1} (z_i - \bar{z})^2} \quad (6)$$

C The Franke function

To study our regression methods we will be using the Franke function, which is a two dimensional weighted sum of four exponentials:

$$\begin{aligned} f(x, y) = & \frac{3}{4} \exp \left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4} \right) \\ & + \frac{3}{4} \exp \left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)^2}{10} \right) \\ & + \frac{1}{2} \exp \left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4} \right) \\ & - \frac{1}{5} \exp \left(-(9x-4)^2 - (9y-7)^2 \right) \end{aligned} \quad (7)$$

D Bias-variance trade-off

We will study bias-variance trade-off by using **bootstrap** resampling technique. [write the theory...]

III Methods

Our first step is to study how the MSE and R^2 changes as the model complexity changes using our three regression methods. For this we will vary

the complexity by the maximum polynomial degree p in our design matrix \mathbf{X} . We will be using the Franke function (7) with $x_i, y_i \in [0, 1]$ while simulated a stochastic noise from a normal distribution $\epsilon \sim N(0, 1)$. We will split our data into a training set and a testing set using **SciKit-Learn**'s `model_selection.train_test_split` function.

Firstly we do OLS regression using this approach with writing our own code from (3), then we do a ridge regression (4) with varying λ -parameter to find an approach that gives us the lowest error. Thirdly we will use **SciKit-Learn**'s implementation of Lasso regression to do the same analysis; that is varying the polynomial degrees.

The next step will be implementing the bootstrap resampling technique using **SciKit-Learn** to do a bias-variance trade-off analysis of the OLS regression. Our goal is to recreate figure 1 (figure 2.11 from [Hastie et al. \(2009\)](#)) which shows an example of idealized overfitting with clear regions of low vs high bias and variance.




Figure 1: Figure 2.11 of [Hastie et al.](#) This idealized graph shows the bias-variance trade-off for our models, and shows the concept of overfitting the training data leading to worse prediction capabilities.

IV Results

test ([Hjorth-Jensen, 2023](#))

[In part a we need to mention that adding stochastic noise increases the MSE, and varies the R2 both negatively sometimes and positively other times (random).]

[in part b we should: "plot some selected results for a specific polynomial degree to illustrate the behavior of the MSE as function of the hyperparameter λ . For the final results however, since you will run many such cases, you should find (search in your table of values) the optimal λ for each polynomial degree and simply tabulate that. Else you will have zillions of figures".]

V Discussion

VI Conclusion

Appendix A Github repository

<https://github.com/LassePladsen/FYS-STK3155-projects/tree/main/project1>

Appendix B List of source code

Here is a list of the code we have developed in this project which can be found in the above Github repository:

- ...
- ...
- ...

Appendix C Analytical derivations

1 Expectation value of y

We will show that $\mathbb{E}(y_i) = \mathbf{X}_{i,*}\boldsymbol{\beta}$ by using $\mathbf{y} = \mathbf{f} + \boldsymbol{\epsilon} \simeq \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$ and separating the expectation value of a sum. Here we approximated \mathbf{f} with $\mathbf{X}\boldsymbol{\beta}$

using OLS. Then taking a value of \mathbf{y} with index i we get $y_i = \sum_j X_{ij}\beta_j + \epsilon_i$:

$$\begin{aligned}\mathbb{E}(y_i) &= \mathbb{E}(\sum_j X_{ij}\beta_j + \epsilon_i) \\ &= \mathbb{E}(\sum_j X_{ij}\beta_j) + \mathbb{E}(\epsilon_i) \\ &= \mathbb{E}(\sum_j X_{ij}\beta_j) \\ &= \sum_j X_{ij}\beta_j \\ &= \mathbf{X}_{i,*}\boldsymbol{\beta}\end{aligned}$$

2 Variance of y

Using the same method as above we will now show $\text{Var}(y_i) = \sigma^2$ where σ^2 is the variance of our data's stochastic noise ϵ . Here we use the definition of variance being $\text{Var}(x) = \mathbb{E}(x^2) - \mathbb{E}(x)^2$:

$$\begin{aligned}\text{Var}(y_i) &= \mathbb{E}(y_i^2) - \mathbb{E}(y_i)^2 \\ &= \mathbb{E}[(\mathbf{X}_{i,*}\boldsymbol{\beta} + \epsilon_i)^2] - (\mathbf{X}_{i,*}\boldsymbol{\beta})^2 \\ &= \mathbb{E}[(\mathbf{X}_{i,*}\boldsymbol{\beta}^2 + \epsilon_i^2 + 2\mathbf{X}_{i,*}\boldsymbol{\beta}\epsilon_i) - (\mathbf{X}_{i,*}\boldsymbol{\beta})^2] \\ &= \mathbb{E}[(\mathbf{X}_{i,*}\boldsymbol{\beta})^2] + \mathbb{E}[\epsilon_i^2] + \mathbb{E}[2\mathbf{X}_{i,*}\boldsymbol{\beta}\epsilon_i] - (\mathbf{X}_{i,*}\boldsymbol{\beta})^2 \\ &= (\mathbf{X}_{i,*}\boldsymbol{\beta})^2 + \mathbb{E}[\epsilon_i^2] + 2\mathbf{X}_{i,*}\boldsymbol{\beta}\mathbb{E}[\epsilon_i] - (\mathbf{X}_{i,*}\boldsymbol{\beta})^2 \\ &= \mathbb{E}[\epsilon_i^2] \\ &= \text{Var}(\epsilon_i) + \mathbb{E}(\epsilon)^2 \\ &= \text{Var}(\epsilon_i) \\ &= \sigma^2\end{aligned}$$

3 OLS expectation value of $\boldsymbol{\beta}_{OLS}$

Here we will show that the expectation value for the optimal $\boldsymbol{\beta}$ for OLS, $\hat{\boldsymbol{\beta}}_{OLS}$, equals $\boldsymbol{\beta}_{OLS}$:

$$\begin{aligned}\mathbb{E}(\hat{\boldsymbol{\beta}}_{OLS}) &= \mathbb{E}[(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}] \\ &= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbb{E}[\mathbf{y}] \\ &= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{X}\boldsymbol{\beta}_{OLS} \\ &= \boldsymbol{\beta}_{OLS}\end{aligned}$$

Here we used that the expectation value of the non-stochastic matrix is just the matrix itself ($\mathbb{E}(\mathbf{X}) = \mathbf{X}$) since it has zero variance (non-stochastic).

4 Variance of β_{OLS}

Here we will show $Var(\hat{\beta}_{OLS}) = \sigma^2(\mathbf{X}^T \mathbf{X})^{-1}$:

$$\begin{aligned}
 Var(\hat{\beta}_{OLS}) &= \mathbb{E}[(\hat{\beta}_{OLS})^2] - \mathbb{E}[\hat{\beta}_{OLS}]^2 \\
 &= \mathbb{E}[(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y})^2] - \beta^2 \\
 &= \mathbb{E}[(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y})(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y})^T] - \beta^T \beta \\
 &= \mathbb{E}[(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \mathbf{y}^T \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1}] - \beta^T \beta \\
 &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbb{E}[\mathbf{y} \mathbf{y}^T] \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} - \beta^T \beta
 \end{aligned}$$

Here we will calculate $\mathbb{E}[\mathbf{y} \mathbf{y}^T]$ separately for ease:

$$\begin{aligned}
 \mathbb{E}[\mathbf{y} \mathbf{y}^T] &= \mathbb{E}[(\mathbf{X} \beta + \epsilon)(\mathbf{X} \beta + \epsilon)^T] \\
 &= \mathbb{E}[(\mathbf{X} \beta \beta^T \mathbf{X}^T + \epsilon \epsilon^T + \mathbf{X} \beta \epsilon^T + \epsilon \beta^T \mathbf{X}^T)] \\
 &= \mathbf{X} \beta \beta^T \mathbf{X}^T + \mathbb{E}[\epsilon \epsilon^T] + \mathbf{X} \beta \mathbb{E}[\epsilon^T] + \mathbb{E}[\epsilon] \beta^T \mathbf{X}^T \\
 &= \mathbf{X} \beta \beta^T \mathbf{X}^T + \mathbb{E}[\epsilon \epsilon^T] \\
 &= \mathbf{X} \beta \beta^T \mathbf{X}^T + \sigma^2 \mathbf{I}
 \end{aligned}$$

Now we put this back into the Variance expression:

$$\begin{aligned}
 Var(\hat{\beta}_{OLS}) &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{X} \beta \beta^T \mathbf{X}^T + \sigma^2 \mathbf{I}) \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} - \beta^T \beta \\
 &= [\beta \beta^T \mathbf{X}^T + (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \sigma^2] \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} - \beta^T \beta \\
 &= [\beta \beta^T + \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1}] - \beta^T \beta \\
 &= \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1}
 \end{aligned}$$

References

- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The elements of statistical learning: Data mining, inference and prediction*. Springer Verlag, Berlin.
- Hjorth-Jensen, M. (2023). *Project 1 on machine learning*. https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/project1.html. ([Online; accessed 26-September-2023])